

FULL LEGAL NAME	LOCATION (COUNTRY)	EMAIL ADDRESS	MARK X FOR ANY NON-CONTRIBUTING MEMBER
Alejandro Piedra Alvarado	Costa Rica	1alepiedra7@gmail.com	
Peter Esekhaigbe	Nigeria	petersonese304@gmail.com	
Jude Chukwuebuka Ugwuoke	United States	jcu0005@auburn.edu	

Statement of integrity: By typing the names of all group members in the text boxes below, you confirm that the assignment submitted is original work produced by the group (excluding any non-contributing members identified with an “X” above).

Team member 1	Alejandro Piedra Alvarado
Team member 2	Peter Eronmosele Esekhaigbe
Team member 3	Jude Chukwuebuka Ugwuoke

Scenario

The team did so well on the previous project that the portfolio strategists want more! The good news is that you learned more: Your toolkit now includes boosting (among other forms of ensemble learning), Support Vector Machines, and Neural Networks. Admittedly, these are more technical and advanced than the previous methods. Nevertheless, as the head strategist (or Spider-Man’s Uncle Ben) said, “With great power comes great responsibility.” Your next responsibility will be to expand the list of best practices for time series. These challenges are more difficult than the first set and include:

**Category 5: Linear
Discriminant Analysis**

**Category 6: Support Vector
Machines**

Category 7: Neural Network

Step 1

Instructor Review for GWP1

Quantitative Analysis (open-ended questions).

- Missing equations in some places.

Writing and Formatting.

- Good organization and writing.

Step 2

- **Category 5: Linear Discriminant Analysis**

- **Basics:**

- Definition: Linear Discriminant Analysis (LDA) is a supervised machine learning algorithm used for classification and dimensionality reduction. It maps data with many dimensions onto a space with fewer dimensions. At the same time, it makes sure different classes are as separate as possible. To do this, LDA finds a mix of features that best tell apart different classes. This mix is linear, meaning it combines features in a straight-line way.
- Classification:
 - **Type:** Supervised Learning
 - **Task:** Classification
 - **Common Use Cases in Finance:** Portfolio risk classification, credit scoring, algorithmic trading signal classification

- **Keywords:** “tags” that identify this model.

- **Supervised Learning:** LDA is a classification algorithm that requires labeled data.
- **Classification:** It assigns data points to predefined categories (e.g., Bullish vs. Bearish markets).
- **Dimensionality Reduction:** Projects data onto fewer dimensions while preserving class separability.
- **Credit Scoring:** Frequently applied in financial risk assessment to classify loan applicants.
- **Financial Risk Modeling:** Helps in identifying financial distress and risk exposure.
- **Market Regime Classification:** Used to distinguish between different market states (e.g., Bull/Bear markets).
- **Portfolio Risk Analysis:** Supports asset allocation strategies by identifying risk regimes
- **Linear Decision Boundaries:** Assumes class separation using linear hyperplanes, simplifying interpretability.
- **High-Dimensional Data Processing:** Effectively handles datasets with multiple financial indicators

- **Category 6: Support Vector Machines**

- **Basics:**

- Definition: Support Vector Machine (SVM) is a machine learning model that uses optimal linear hyperplanes to separate data points according to predefined classes and

labels. It is generally applied for solving complex classification and regression problems. The positioning and dimensional space of the separating hyperplane can be adjusted (tuned) by the modifying hyperparameters. The hyperplane is usually bound by two lines called *support vectors*, which help to maximize the margin between two different classes on either side of the linear hyperplane. The shorter the perpendicular distance of the support vectors is known as the margin, and an optimal hyperplane is obtained when the margins are maximised.

While there are a number of other hyperparameters tuned to adjust the hyperplane, especially when dealing with non-linear data, the essence of SVM is primarily to reduce misclassification and margin errors in classification data points

- **Classification:** SVM is classified based on the type of learning, classification problem, and dimensionality of the functions
 - Type: Supervised learning
 - Subcategory:
 - A. Based on the learning type, SVM is classified into two categories:
 - a. Linear SVM: which applies data points that can be separated by a linear hyperplane.
 - b. Non-linear SVM: this applies transformation (kernel) functions for separating non-linear data points.
 - B. Based on the classification problem, there are problem types: SVM can be a *binary classification SVM split* (classifies data into two categories), a *multi-class SVM splits* (classifies data into multiple categories using different approaches), a *regression SVM* (useful for continuous values), or a *one-class SVM* (which is useful for identifying patterns).
 - C. Based on the kernel functions: SVM can be categorised into *linear SVM* (applies simple dot product in the linear equation), *polynomial SVM* (involves polynomial transformations for variables in the equation), *radial basis SVM function* (deal with high-dimensional hyperplane spaces for non-linear data), and *neural network SVM functions* (such as sigmoid function)
- **Keywords:** Hyperplane, classifiers, soft margins, hard margins, regularization parameter, hyperparameter, SVM, kernel functions, maximal margin classifier, SVM regressor

- **Category 7: Neural Networks**

- **Basics:**
 - **Definition:** A Neural Network (NN) is a computational model inspired by the human brain, designed to recognize data patterns. It consists of layers of interconnected nodes

(also called neurons), where each connection carries a weight that adjusts during training. These models can approximate highly non-linear functions, making them universal function approximators. The networks are trained as if the whole brain acts as a single neuron, thus making themselves the universal, non-linear function approximators.

■ Classification:

Type	Description
Feedforward NN	Data flows from the input to the output without any cycles in the simplest form.
Multilayer Perceptron (MLP)	It utilizes a single or many delayed layers, which is currently the most popular model used in finance.
Convolutional NN (CNN)	Not only is it specifically intended for two-dimensional grid-like data, but people are also declining to use it in finance while they use it more in image processing.
Recurrent NN (RNN)	Networks with loops that can be more useful for time-series and sequence modeling in financial data.
Deep Neural Network (DNN)	Neural Networks that are constructed of multiple hidden layers applicable to tasks such as structuring complex data usually post real-world data usage.

How It Works

- Input layer: It accepts unprocessed data, e.g., stock prices and macroeconomic indicators.
- Hidden layers process at least one hidden layer where the inputs are transformed using weighted sums and activation functions: ReLU, sigmoid, or tanh.
- Output layer: Produces prediction results (e.g., “Buy” or “Sell” signal)

Each neuron calculates:

$$z = \sum_{i=1}^n \omega_i x_i + b$$

And then it applies an activation function to generate the neuron's output.

Key Characteristics

- It learns through backpropagation, which aids in reducing a loss function, which can be (for example) the cross-entropy function for classification.
- It has been improved through the algorithm of stochastic gradient descent (SGD) and Adam.
- Solve complex and non-linear relationships in high-dimensional datasets.

Application in Finance

- Market forecasting: Analyzing the future of stocks, bonds, or commodities to use them as decisive factors in establishing market positions.
 - Portfolio optimization: Assigning weights to assets depending on recognized connections.
 - Credit risk modeling: Evaluating the likelihood of default or predicting credit
 - Sentiment analysis: Categorizing news or social media sentiment concerning financial assets.
- **Keywords:** “tags” that identify this model.
 - #Neural Networks, #Multilayer Perceptron (MLP), #Bridged MLP (BMLP), #Deep Learning, #Levenberg-Marquardt, #Backpropagation, #Financial Forecasting, #Nonlinear Modeling, #Time Series Classification, #Signal Processing

Step 3

Category 5: Linear Discriminant Analysis

Advantages:

- **Class Separability:** Makes the difference between financial groups (like credit risk classes) as big as possible
- **Dimensionality Reduction:** Useful for high-dimensional financial datasets while retaining key information.
- **Interpretability:** Unlike more complex non-linear models, LDA shows clear decision boundaries.
- **Robustness:** Works well even with small sample sizes.
- **Computationally Efficient:** Faster than complex non-linear classifiers like Support Vector Machines and Neural Networks.

Computation:

1. **Compute Class Means:**
 - The code defines the target variable (data['Target']), with classes representing bullish (1) and bearish (0) markets.
 - The Linear Discriminant Analysis (LDA) model (LinearDiscriminantAnalysis(solver='eigen', shrink='auto')) inherently figures out group averages.
2. **Compute Scatter Matrices:**
 - The LDA model guesses how spread out things are within groups, meaning it estimates the within-class scatter matrix S_W and between-class scatter matrix S_B .
 - The eigenvalue math problem is solved via the solver='eigen' parameter.

3. **Compute Discriminant Function:** The `lda.fit(X_train, y_train)` step computes the transformation that maximizes class separability by solving $S_W^{-1} S_B$.
 4. **Project Data:** The `X_lda = lda.transform(X_test)` step puts test data on a line that shows group differences best.
 5. **Classify New Data:**
 - The `lda.predict(X_test)` function assigns new observations to the most probable group or class by figuring out how each group is with `lda.predict_proba(X_test)`.
- For more details on the computational implementation and the overall process on how this works, refer to “MLiF_GWP2_g8507.ipynb”, Category 5: Linear Discriminant Analysis.

Disadvantages:

- **Assumes Normality:** LDA assumes that data is normally distributed, which financial datasets often fail to meet.
- **Linear Boundaries:** Struggles to handle complex decision boundaries in financial classification problems that aren't linear.
- **Sensitive to Outliers:** Financial anomalies can throw off how well the classification works.
- **Equal Covariance Assumption:** Assumes that all classes share the same covariance structure, which is rarely the case in financial markets.
- **Limited Performance in High Dimensions:** Struggles with datasets where the number of features exceeds the number of observations, leading to instability in matrix inversions.

Equations:

- Within-class Scatter Matrix: sums up the covariance matrices of every class to measure how much the data is spread out in each one.

$$S_W = \sum_c \sum_{i \in c} (x_i - \mu_c)(x_i - \mu_c)^T$$

- Between-class Scatter Matrix: shows how much the average values differ from one class to another.

$$S_B = \sum_c^c N_c (\mu_c - \mu)(\mu_c - \mu)^T$$

- Eigenvalue Problem: main challenge when working with LDA, where solving for eigenvalues (λ) and eigenvectors (v) determines the best direction to project the data, making sure different classes can be told apart.

$$S_W^{-1} S_B v = \lambda v$$

- Discriminant Function: Projects data onto the new axis and assigns class labels based on:

$$y = w^T x + w_0$$

Where $w = S_W^{-1}(\mu_1 - \mu_2)$ is the projection vector and w_0 the threshold for classification.

Features:

- Works with categorical financial data (e.g., credit risk classifications).
- Handles high-dimensional datasets efficiently while preserving critical financial insights.
- Provides class probability estimates, useful for risk modeling.
- Assumes normally distributed predictors, which makes sense and is often reasonable for financial metrics after some transformation.
- Utilizes linear decision boundaries, making it easier to understand for regulatory and risk management applications.
- Can be extended to quadratic discriminant analysis (QDA) for more complex decisions.
- Often used as a benchmark for more complex machine learning models in financial classification problems.

Guide:

- Inputs

- Market-Based Features (From Yahoo Finance):
 - **S&P 500 (^GSPC)**: Main market indicator (*used for labels, not features*)
 - **Nasdaq-100 (^NDX)**: Captures tech-sector trends.
 - **VIX (^VIX)**: Measures market fear and uncertainty
 - **US Treasury 10-year yield (^TNX)**: Reflects long-term interest rate expectations.
 - **Gold (GC=F)**: Risk-off asset, often inversely correlated with equities.
 - **Bitcoin (BTC-USD)**: Captures speculative risk-on sentiment.
- Macroeconomic & Credit Indicators (From FRED).
 - **Federal Funds Rate (FFR)**: Fed's monetary policy stance
 - **Corporate Bond Spread (Baa - Aaa Yield)**: Measures credit risk in corporate bonds.
- Data Processing Steps
 - **Daily Returns Calculation**: Convert raw prices into daily log returns for stationarity.
- Labels (Target Variable)
 - **Market Regime Classification**:
 - **Bullish (1)**: S&P 500 daily return > 0.
 - **Bearish (0)**: S&P 500 daily return ≤ 0.

- Outputs

- **Class Predictions**: Assigns each observation into the most probable financial group (e.g., Bullish/Bearish market conditions).

- **Class Probability Estimates:** Computes posterior probabilities, which help to assess risk and weight portfolios.
- **Feature Importance:** Gives coefficients to each financial indicator, showing their impact on classification.
- **Linear Discriminants:** Produces transformed feature representations that maximize class separability.
- **Dimensionality Reduction:** Maps high-dimensional data to an optimized space with fewer dimensions.
- **Performance Metrics:**
 - **Confusion Matrix:** Measures correct vs. incorrect classifications.
 - **Classification Report:** Displays accuracy, completeness, and overall score.
 - **Accuracy Score:** Provides the model's overall performance measure.
- **Financial Decision Support:**
 - **Market Trend Analysis:** Identify bullish or bearish trends for trading strategies.
 - **Portfolio Optimization:** Adjust asset allocation based on market regime classification.
 - **Credit Risk Assessment:** Evaluating borrower default risk using financial indicators.

Hyperparameters:

- **Solver:** This parameter determines the algorithm used for optimization.
 - 'svd' (Singular Value Decomposition):
 - Default, generally preferred for datasets with many features. Does not support shrinkage.
 - 'lsqr' (Least Squares):
 - Can be faster for large datasets, but may require data to be centered.
 - 'eigen' (Eigenvalue Decomposition):
 - Useful for datasets with a small number of features.
- **Shrinkage / Regularization Parameter:** Helps when S_w is singular or poorly conditioned.
 - None (default): No shrinkage.
 - 'auto': Automatic shrinkage estimation using the Ledoit-Wolf lemma.
- **Number of Components:** Controls the number of discriminant axes (components) used for dimensionality reduction. For cases involving binary classification, changing this value from 1 won't make a difference.
- **Priors:** Allows you to set the class priors.
 - If not set, the class priors are determined by the training data.
 - Useful if one knows the actual distribution of your target variable.

Illustration:

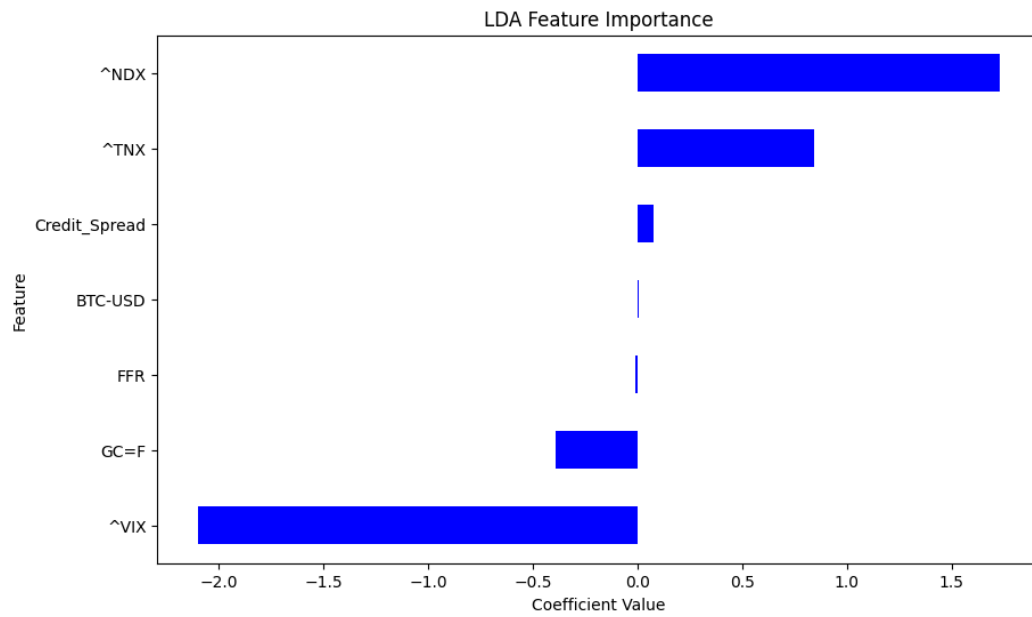


Figure 1. LDA Feature Importance.




Figure 2. LDA Decision Boundary.




Figure 3 (a). LDA: Bullish vs. Bearish (Training Data).

Figure 3(b). LDA: Bullish vs. Bearish (Testing Data).




Figure 4. Confusion Matrix.

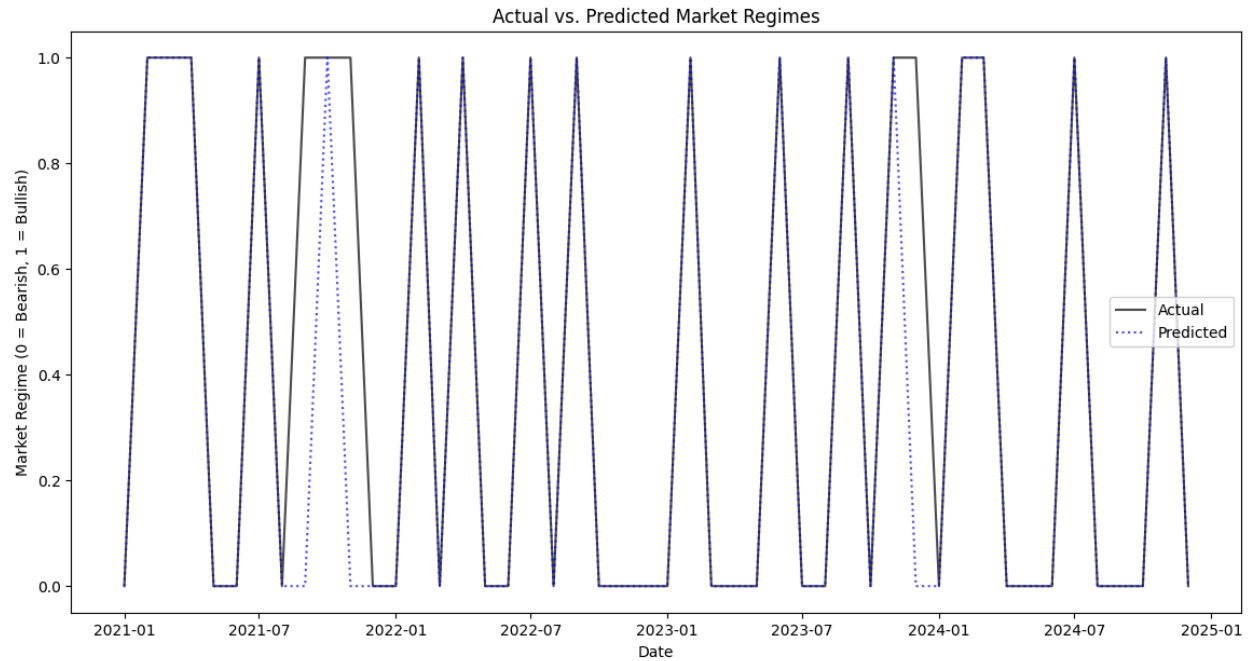


Figure 5. Actual vs Predicted Market Regimes.




Figure 6 (a). LDA Bullish Market Probability Distribution.




Figure 6 (b). LDA Bearish Market Probability Distribution.




Figure 7. Market Regime Classification Over Time.

Journal:

- Alrasheedi, Melfi. "Predicting Up/down Direction Using Linear Discriminant Analysis and Logit Model: The Case of SABIC Price Index." *Research Journal of Business Management*, vol. 6, no. 4, Apr. 2012, pp. 121–33, <https://doi.org/10.3923/rjbm.2012.121.133>.

This paper applies Linear Discriminant Analysis (LDA) alongside the Logit model to predict the up/down direction of the SABIC Price Index, a key financial asset. By combining both technical and fundamental features, Alrasheedi aims to enhance market forecasting, offering an approach that classifies price movements into bullish and bearish categories. The study leverages historical price data and market indicators to identify market trends and assist investors in making more informed decisions. The paper gives a good starting point for using LDA to predict price changes in finance, which is key to managing portfolios and assessing risks. The results are relevant for those managing assets and traders looking for systematic tools to predict how markets will act and change how they invest based on that.

Relevance to Our Project:

This paper applies LDA to classify market trends, which matches our project's approach. We also use LDA to categorize bullish and bearish market conditions using various financial indicators. The study combines technical and fundamental data for classification just like our method. We also use market data such as S&P 500, VIX, and other financial metrics to predict market movements. Alrasheedi's research confirms that LDA works well for managing portfolios and making financial decisions. This suggests we could adapt it for our model too.

- **Category 6: Support Vector Machines**

Advantages: Some of the advantages of Support Vector Machines are as follows:

- **Flexibility in handling Non-linear Data:** Support vector machines are effective for classifying non-linear data with kernel functions. The flexibility of non-linear kernel functions in finding suitable decision boundaries for non-linear data makes the support vector classifier an appropriate algorithm for handling non-linear data.
- **Effective for Small Datasets and High-dimensional Spaces :** Unlike other supervised learning algorithms like KNN and decision trees that may struggle with handling very high-dimensional spaces, the SVM algorithm performs well with high-dimensional data with a large number of dimensions that exceed the number of available samples. Additionally, SVM also works well with small structured datasets.
- **Maximum Margin Classifiers:** SVM optimises margin between classes: perpendicular distance between support vectors and the optimal hyperplane can be adjusted to maximise distance between data points in different classes. This in turn minimises misclassification errors and margin violations and improves generalisation.
- **Generalization:** Support Vector Machines algorithm is robust enough to reduce overfitting of the training data yet able to control the trade-off between complexity and classification accuracy.
- **Versatility:** SVM is versatile as it applies to both classification and regression problems. While support vector classifiers are commonly used for classification problems, support vector regressors are used for regression problems.

Computation: Refer to “MLiF_GWP2_g8507.ipynb” for the computational implementation details.

To illustrate the Support Vector Machine Algorithm, we will consider a loan prediction dataset sourced from Kaggle. The goal of the dataset is predict whether a loan applicant's loan application will be approved or rejected. The dataset contains 13 different loan features (and 614 samples) ranging from the borrower's gender, marital status, educational background, loan ID, loan amount, credit history, etc. For the dataset, we will implement different approaches:

- **Approach 1:** Apply the SVM classifier to split data using Loan status as the target class. Additionally, we'll determine the optimal tuning hyperparameters to use in the classification problem.
- **Approach 2:** Apply SVM Regressor to model the data comparing Predicted Loan Amount with Actual Loan Amount. Additionally, we will also compare the regressor model's performance with that of other regression models.

A. For Approach 1: SVM Classification Task:

- **Load Packages:** Loaded in the relevant Python Libraries as well as the dataset “loan_predictor.csv” data
- **Exploratory Data Analysis and Data Preprocessing:** We proceeded to implement exploratory data analysis and relevant data preprocessing, highlighting relevant variables and removing inconsistent NaN values.

- **Model Training:**

- We identified *Loan_Status* as the target variable, where *Loan_Status* value of 1 signifies that applicant is eligible for loan, while value of 0 indicates the applicant is NOT eligible for a loan
- More importantly, we categorised and visualised relevant features in different seaborn plots based on *Loan_Status*.
- Furthermore, we removed columns (attributes) with NaN values before running the SVM classifier model. This is because NaN values hamper full hyperparameter selection implementation.
- Specify *Loan_Status* as the y (target) variable and other columns as X variables
- Implement a train-test set split in ratio 80%-20% split, respectively.
- Run a grid search using the *GridSearchCV* cross validator to determine the best kernel and slack variable (regularization parameter, C), over 5 validations. The grid search ran over a list of kernels: "*linear*", "*poly*", "*rbf*." The C parameters range list: 0.01, 0.1, 1, 10, 100. Upon running the grid search, *SVC* was selected as the best estimator, *poly* selected as the best kernel function, and 0.01 and the most appropriate regularization parameter C.
- Train the data with the *tuned* SVM classifier (i.e., the SVM classifier with the best hyperparameters). We did the same with the decision tree classifier to compare the performance of the SVM classifier after all.

- **Evaluate Model Performance:**

- We compared the performance of the SVM classifier, DecisionTreeClassifier and Random guess on an ROC curve. Apparently, the SVM classifier performs better than the random guess, and its AUC result is pretty similar to that of a DecisionTree classifier.
- Further, we evaluated the performance of the SVM classifier on the test set; we obtained an accuracy score of 65%, which is a good enough score for a machine learning model. Lastly, we detailed the performance of the model in a classification report table

B. Approach 2: SVM Regression Task:

- Load packages and original dataset ("*loan_predictor.csv*")
- Implement data preprocessing.
- Specify *LoanAmount* as the y (target) variable and the remaining variables as X variables.
- Implement *Train-Test split* as in the similar ratio as the SVM classifier
- Normalise the train and test set using scikit-learn's *StandardScaler()* for proper feature scaling.
- Run the SVM Regression Model:
 - We make use of the radial basis kernel ("*rbf*") function instead of the polynomial ("*poly*") kernel as in the SVM classifier because the former is more suited for complex non-linear relationships.

- Having specified the kernel function, we run a grid search to determine the optimal values for other hyperparameters such as *regularisation parameter C*, *gamma*, and *epsilon* (since we'll be working with a regressor). From the results, $C = 100$, $\gamma = 0.01$, and $\epsilon = 0.001$ were obtained
- Train the data with a tuned model (with the *best hyperparameters*).
- Evaluate the performance of the SVM regressor using *Mean Squared Error (MSE)*. We obtained an MSE value of 2572.07 for the SVM Regressor.
- Visualise the SVM Regressor Model:
 - First we extract the supporting vectors from the X_{train} and y_{train} set.
 - Plot the SVM Regressor Model in a scatter plot. Add a *Perfect Prediction* regression line to the plot as well as the supporting vectors showing the margin tolerance between data points on both sides of the regression line.
- Compare SVM Regressor Model with other Regression Models:
 - First, we ran two different regression models: Random Forest and XGBoost models, and calculated their respective MSE values.
 - Upon comparison of the 3 models, SVM regression model had the least MSE, hence the most appropriate model for the dataset.

Disadvantages: Some of the disadvantages of the Support Vector Machines algorithm are as follows:

- **Computational Efficiency:** Due to its application to high-dimensional but large datasets, the SVM algorithm may become computationally expensive when working with complex kernel functions
- **Scalability:** While SVM may perform well on datasets with many attributes (high-dimensional data), it may struggle with significantly large-scale data. Therefore, SVM is effective for classifying only small-scale and medium-scale data.
- **Challenging Hyperparameter Tuning:** Adjusting hyperparameters such as the regularization term C , selecting the right non-linear kernel functions could be challenging and therefore requires careful consideration. No kernel function can generalise all situations, and the number of available kernels may be limited. In the Jupyter notebook example, "linear" performed slower than other kernel functions when deciding on the best estimators to use for the model. In fact, as the regularisation parameter value increased, computational (iteration) time increased as well.
- **Outliers Sensitivity:** SVM is sensitive to noisy data and mislabeled data. This consequently can affect the performance of the algorithm.
- **Interpretability:** The outputs of SVM algorithms are not as interpretable as those of decision trees and linear regressions since they are not probabilistic. Hence, they require additional methods to obtain probabilistic scores for classes.

Equations:

The equations that summarise the SVM model are dependent on the kernel function being used for modelling. Mathematically, the SVM is represented by the equation of the hyperplane (separating boundary) between classes.

Hence, we categorise the following equation based on the kernel function:

- **Linear Function Kernel:** For a linear support vector kernel classifier, the hyperplane is represented mathematically as:

$$f(x) = \beta_o + \sum_{i=1}^n \alpha_i(x, x_i)$$

where:

$f(x)$ = Linear function

β_o = Intercept

α_i = Coefficients for each i component

x, x_i = two observations in the datasets

The product of observations can be represented as a general kernel function. Hence, a linear kernel function can be rewritten as:

$$K(x_i, x_{i'}) = \sum_{j=1}^p x_{ij} x_{i'j}$$

where:

$K(x_i, x_{i'})$ = Linear kernel function (quantifies the similarity of two observations)

The linear kernel quantifies the similarity of observations using Pearson correlations

- **Polynomial Kernel Function:** The polynomial function is used to show the relationship between the moderate non-linear data. Mathematically, it is represented as:

$$K(x_i, x_{i'}) = (1 + \sum_{j=1}^p x_{ij} x_{i'j})^d$$

where:

$K(x_i, x_{i'})$ = Polynomial kernel function (quantifies the similarity of two observations)

d = Degree (> 1)

- **Radial Kernel Function:** The radial kernel function is used to show the relationship between the complex non-linear data variables. Mathematically, it is represented as:

$$K(x_i, x_{i'}) = \exp(-\gamma \sum_{j=1}^p (x_{ij} - x_{i'j})^2)$$

where:

$K(x_i, x_{i'})$ = Radial kernel function (quantifies the similarity of two observations)

γ = gamma (which measure the influence of single training example)

$(x_{ij} - x_{i'j})^2$ = Euclidean distance between data points in each observation

While the aforementioned kernel functions are the most commonly used ones, there are other ones, like the sigmoid function, not considered in this report.

- **Maximum Margin Classifier:** The Margin Classifier measures the margin tolerance between the support vectors and the optimal hyperplane. They acts as buffers against misclassification errors and margin violations. The maximal margin classifier maximises the distance of the optimal hyperplane to the nearest data points. Mathematically, the maximal margin classifier is subject to:

$$\sum_{j=1}^p \beta_j^2 = 1$$

Such that:

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M \quad \forall i = 1, 2, \dots, n$$

Maximal margin classifiers are influenced by *slack variables* such as regularization parameter C and epsilon ϵ such that we:

Maximizing M is subject to

$$\sum_{j=1}^p \beta_j^2 = 1$$

Such that:

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i) \quad \forall i = 1, 2, \dots, n$$

$$\epsilon_i \geq 0, \quad \sum_{j=1}^n \epsilon_i \leq C$$

$\epsilon_i = \text{epsilon (measures the margin tolerance)}$

Features:

- **Works well with high-dimensional spaces:** SVM is effective in high-dimensional spaces, even when the number of features is greater than the number of samples.
- **Versatility in Handling Classification and Regression Problems:** While the SVM classifier helps separate data into classes, the SVM regressor helps with predicting continuous values. In the Jupyter notebook example, we use the SVM classifier to classify the data set based on the target variable *Loan Status* while *Loan Amount* was predicted using the SVM Regressor.
- **Ideal for Small and Medium-scale Data:** While the SVM classifier can handle as many attributes as possible, it cannot handle large-scale data with large number of rows. It is suitable for small and medium-scale data.
- **Handles Non-linear data relationships:** SVM can handle both linear and non-linear separable data well. Kernel functions help in data transformations so as to allow for as easy separations.
- **Permits Hyperparameter Tuning:** Slack variables such as the regularisation terms C , γ , and ϵ help kernel functions maximise the margin classifiers. More importantly, this helps to balance model complexity and accuracy.
- **Robust to Overfitting:** Tuning the regularisation parameter C helps mitigate against overfitting. High C values run the risk of overfitting the model to the training set and vice versa.

Guide:

Inputs

- Data source on which SVM classification or regression is to be implemented.
- Specified X (features) variables and y (target variables): For a classification problem, the y variable serves as the basis for classification of the rest of the data. For a regression problem, prediction of continuous variables is implemented on the y-variable.
- Train-test split: The train-test split ratio must also be defined. Traditionally, the proportion of the train set is always larger than that of the test set so as to allow for the machine learning model to learn adequately from the more training data points.
- Hyperparameters: The best tuning parameters (C , γ , and ϵ) is selected for model tuning so as to obtain optimal hyperplane, minimal classification errors, and margin violations. Additionally, we could also determine the best kernel function to use for model training using grid search with a number of cross validations.
- Model development: The tuned model (SVM classifier or SVM Regression) is used to train the data. Apply the tuned model to the test set afterwards.

Outputs

- Optimal Hyperplane/Regression Line: For a classification problem, we obtain an hyperplane separating categorical target variable into different classes. For a regression problem, we obtain prediction values for the continuous variables in the test set.
- Performance/Accuracy Score: Helps to evaluate the performance of the SVM models. For a SVM classifier, the accuracy of the model can be visualised on a Receiver Operating Curve (ROC) in comparison with a DecisionTree Classifier and a Random guess. For the SVM Regressor, the accuracy can be measured by metrics such as Mean-Squared Error (MSE), R^2 , or RMSE (Root Mean Squared Error). Compare the accuracy of the SVM regressor model with that of other regression models to see if it performs better.

Hyperparameters:

- **Kernel-specific hyperparameters:** The “linear” kernel function can be used for standard linear data. For non-linear data, kernel functions such as: “poly”, “rbf”, “sigmoid” etc. Other kernel-specific functions may include the *degree* of polynomial function.
- **Common hyperparameters:** This consists of the *regularization parameter* C (which regulates the trade-off between maximizing the margin and minimizing classification errors). *Gamma* (γ) controls the influence of data points on the decision boundary; Lower (γ) implies a smoother decision boundary and vice versa. *Epsilon* (ϵ) controls the margin tolerance range.
- **Additional hyperparameters:** These include *cv* and *probability* parameters, which help to refine the results of the trained model. When the *probability* parameter is set to True, the model adopts probability estimates and may consequently result in slower training. *cv* parameter on the other hand stipulates the number of cross validation steps that should be implemented during grid search.

Illustration:

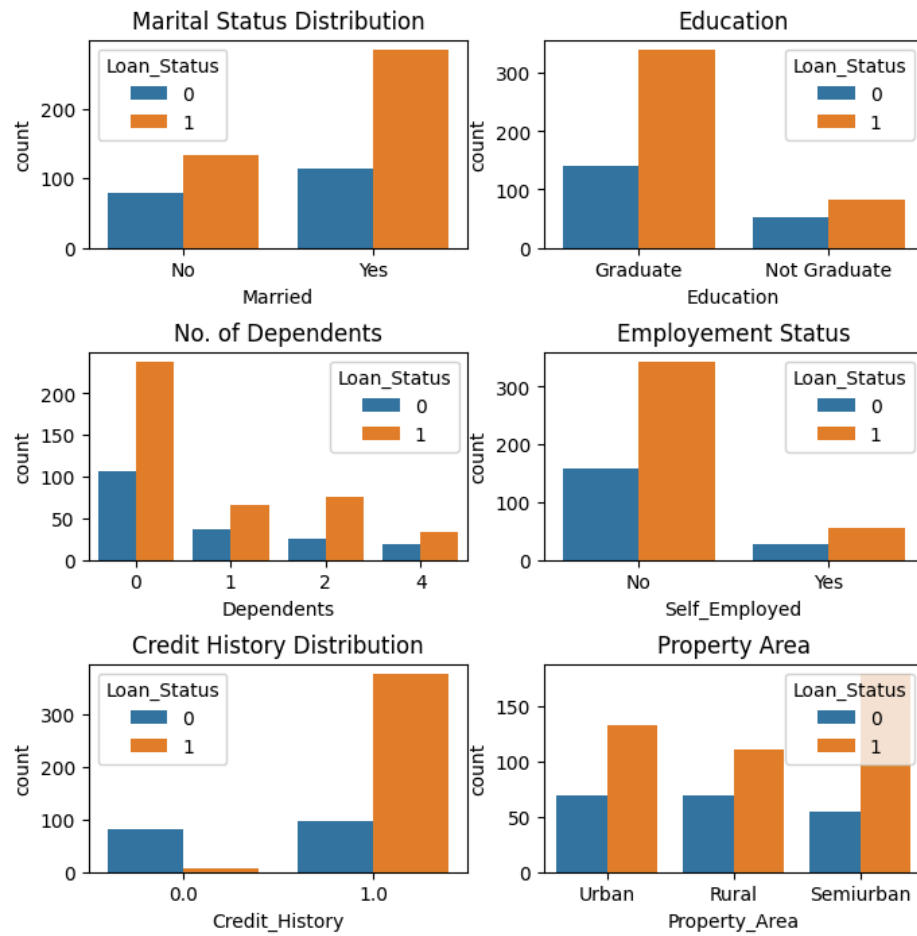


Fig. 8. Features-Target Distribution




Fig. 9. Receiver Operating Curve




Fig. 10. SVM Regressor: Actual vs Predicted




Fig. 11. SVM Regressor Model Performance Comparison

Journal:

- Mahmoud K. Okasha "Using Support Vector Machines in Financial Time Series Forecasting." Department of Applied Statistics, Al-Azhar University, Gaza Palestine, International Journal of Statistics and Applications, p-ISSN: 2168-5193 e-ISSN: 2168-5215, 2014; 4(1): 28-39, doi:10.5923/j.statistics.20140401.03

Description: This article describes the application of support vector machines in stock financial times series forecasting. The complexity of stock price indices and the noisiness of the time series further complicate forecasting. Hence, in this paper, algorithms ARIMA, ANN, as well as SVM, were fitted to Al-Quds Palestinaian Stock Exchange ultimately to forecast two-month further data. The performance of all three models was compared using root-mean-squared error. It was concluded from the paper that SVM was a more accurate model for forecasting. The results from this paper are consistent with the results obtained in the Jupyter notebook example where we compared the performance of SVM regressor models with the regressor models like Random Forest and XGBoost.

Category 7: Neural Networks

Advantages:

- Neural networks can learn and model complex, non-linear relationships in a way that surpasses all conventional algorithms.
- They are also excellent at handling superposed data with many input features simultaneously, sometimes more straightforward statistics.
- Neural networks are capable of adapting to new data by gradually updating weights, which enhances performance over time.

- When set at the proper tuning level, they are capable of effectively generalizing unseen data by implementing techniques such as dropout and regularization.
- Neural networks prefer the parallel computation of independent neurons in a single layer.
- Deep learning networks excel at extracting the necessary features from raw data. In other words, they are useful in automating tasks.
- Image recognition, natural language processing, financial time series forecasting, and medical diagnosis are common practices that widely employ this technology.
- Sufficiently able to cope with noisy or incomplete data makes them perfect for real-world tasks.
- New tasks can use the already trained models, and this can save time and money.
- It is possible to approximate any continuous function by using a sufficiently large neural network.

Computation of Neural Networks

- Overview

This outlines the computational steps for constructing and assessing a Neural Network (MLP Classifier) using financial and macroeconomic information. The principal aim of this study is to anticipate changes in the stock market, primarily concentrating on the performance of the S&P 500 Index (^GSPC).

- Data Preparation

Financial market data from November 2017 to January 2025 was gathered from Yahoo Finance using GSPC, NDX, VIX, TNX, GC=F, BTC-USD, and ETH-USD. Furthermore, the indicators of the US economy, the Federal Funds Rate (FFR) and Baa and Aaa-rated bond yields, were retrieved from FRED.

First, the daily returns from all market data sets were computed. After that, the economic data was processed in the form of a credit spread, i.e. a difference between Baa Yield and Aaa Yield, which is a widely recognized gauge of market risk.

The datasets were then combined systematically by matching the corresponding date in both sets. The target variable was specified as a binary variable where '1' means the S&P 500 index had a positive return on that day while '0' indicates the opposite.

- Model Development

The data was divided into train and test subsets, which contained 80% of the training set and 20% of the testing set, thus keeping the sequence untouched, hence the existence of the chronological order within the time series analysis.

All features were standardized to allow the efficient training of the neural network.

To implement a neural network, the MLPClassifier of the scikit-learn library was used, where the neural net consists of two hidden layers (16 neurons in the first layer and 8 ones in the second layer). The 'ReLU' activation function was used with 'adam' as the optimizer, and the maximum number of iterations was set to 300, which was the training for all of them.

- Evaluation and Results

The model achieved an accurate score of 75%. The detailed results were as follows:

```
Confusion Matrix:
[[4 1]
 [2 5]]

Classification Report:
              precision    recall  f1-score   support

     0       0.67       0.80       0.73         5
     1       0.83       0.71       0.77         7

 accuracy          0.75
 macro avg         0.75       0.76       0.75        12
 weighted avg      0.76       0.75       0.75        12

Accuracy Score: 0.75
```

Through this computational illustration, we demonstrate the high efficiency and robustness of NNs in predicting the movement of financial markets using financial data and macroeconomic data together.

Disadvantages

- Neural networks, lacking in interpretability, are often critiqued.
- Training in deep networks requires extensive resources, including GPUs and considerable memory.
- If neural networks are not properly regularized, they tend to learn or memorize data from the training dataset instead of generating proper and generalizing associations.
- By using more data for training, neural networks tend to get improved performance, whereas, for small datasets, they perform less effectively.
- Several hyperparameters (e.g., learning rate, batch size, architecture) must be optimized carefully.
- Deep networks can be a very time-consuming part of training.
- Poor initialization can lead to vanishing/exploding gradients or slow convergence.
- While uncommon with modern optimizations, local minima can also be encountered in neurons.
- In complex networks, the gradients may disappear or grow excessively during backpropagation.
- Input features must be scaled (e.g., normalized or standardized) to prevent the model from being improperly convergent.

Equations of Neural Network

Name	Equation	Brief Explanation
Weighted Sum at Each Neuron	$z = \sum_{i=1}^n \omega_i x_i + b$	Here, input features are multiplied by weights and added with a bias term for each equation

Activation Functions	These equations below are examples of activation function	Here, we pass z through an activation function to introduce non-linearity
Sigmoid	$\sigma(z) = \frac{1}{1 + e^{-z}}$	This maps any real-valued number into a range between 0 and 1 for binary classification.
ReLU (Rectified Linear Unit)	$\text{ReLU}(z) = \max(0, z)$	This sets negative values to zero and keeps positive values unchanged
Tanh	$\frac{e^z - e^{-z}}{e^z + e^{-z}}$	This outputs values between -1 and 1 that are centered around 0 to make the optimization easier than sigmoid
Forward Propagation (Layer-by-Layer Computation)	$a^L = f(w^L a^{L-1} + b^L)$	Here, each layer computes an output/activation through the application of weights and biases that are followed by non-linear activation function.
Loss Functions (e.g., for Classification)	Loss functions are exemplified below	This measures the difference between the predicted outputs and the true values in order to quantify a model
Binary Cross-Entropy Loss	$L = -m \ln = \frac{1}{m} \sum [y(i) \log(y^{(i)}) + (1 - y(i)) \log(1 - y^{(i)})]$	This is used to penalize incorrect predictions in a binary classification.
Backpropagation & Gradient Descent	The equations are exemplified below	This optimizes weights through the computation of gradients of the loss function with respect to each of the parameters.

Compute gradients	$\frac{\partial L}{\partial w}$ and $\frac{\partial L}{\partial b}$	This determines how much each of the weights or biases contributed to the loss in order to update the parameters.
Update rule	$w = \eta \frac{\partial L}{\partial w}$	This updates the weights by using a learning rate and the computed gradients in order to minimize loss.

Features

- A layered network model has an input layer, a hidden layer, and an output layer. It allows the computer to learn hierarchically.
- It is possible to reproduce any recursive function in Universal Approximator by just using sufficient data and computational powers in the neural network.
- Non-linearity is used to recognize intricate patterns by employing activation functions.
- Backpropagation, which learns through adjusting weights using the loss function's gradients.
- Handles high-dimensional data and can process datasets with large feature spaces effectively.
- Transfer Learning Ready, which can use pre-trained models and fine-tune for new tasks.
- Parallelizable computation that is suited for GPUs and distributed processing.
- Supports regularization which comprises techniques such as dropout and L1/L2 to mitigate overfitting.
- Flexible output Structures that are fit for types of classification, regression, and even multi-output problems.
- It can be architecturally tuned in varying arbitrary depths, widths, and activation functions for different uses or field cases.

Guide

- Inputs
 - Feature Matrix X: This is a numerical matrix of shape (m,n), where m is the number of examples and n is the number of features. Examples are prices, volumes, indicators, or pixel intensities.
 - Target Labels y: This is a vector or matrix of ground truth values (e.g., it has 0/1 for binary classification and categories for multi-class).
 - Hyperparameters: These include learning rate, number of layers, activation functions, optimizer, batch size, epochs, etc.

- Outputs

- Predictions: Classification: probabilities or class labels. The regression is a continuous value.
- Loss Score: This numeric value represents prediction error (e.g., cross-entropy or MSE).
- Model Weights and Biases: These are internal parameters that are learned during training.
- Performance Metrics: These include accuracy, precision, recall, F1-score, AUC, etc.

Hyperparameters of the Neural Network

- Learning Rate (η): Controls the step size during gradient descent updates. This can be too high (divergence) or too low (slow convergence).
- Number of Epochs: Number of complete passes through the training dataset.
- Batch Size: Number of samples that are processed before the model is updated.
- Number of Hidden Layers & Neurons
- Activation Functions
- Optimizer
- Algorithms to update weights: SGD, Adam, RMSProp.
- Dropout Rate
- Weight Initialization
- Regularization Parameters
- Learning Rate Scheduling

Illustrations

Daily Returns




Fig. 12. S&P 500 and NASDAQ Daily Returns

The chart above depicts extreme market oscillations, such as the one experienced in the first quarter of 2020 and brings into focus how both market indices respond almost identically to changes in the macroeconomic climate. NASDAQ sometimes reveals slightly more significant alterations due to its tech-centered structure.

Correlation Heatmap

The heatmap below shows several significant connections: Bitcoin and Ethereum have a remarkable similarity (0.79), meaning that, in most cases, they hold the same direction. The S&P 500 and NASDAQ are twin siblings (0.94) as they share components, whereas the S&P 500 and VIX are mirror images of one another (-0.76), which shows the fear of the market. In the above, the rise of the indices over the last month triggered the downtrend in Bitcoin (-0.30), while the wider credit spreads affecting the S&P 500 were negatively correlated (-0.37), hinting at the financial strain.




Fig. 13. Correlation Heatmap

Credit Spread

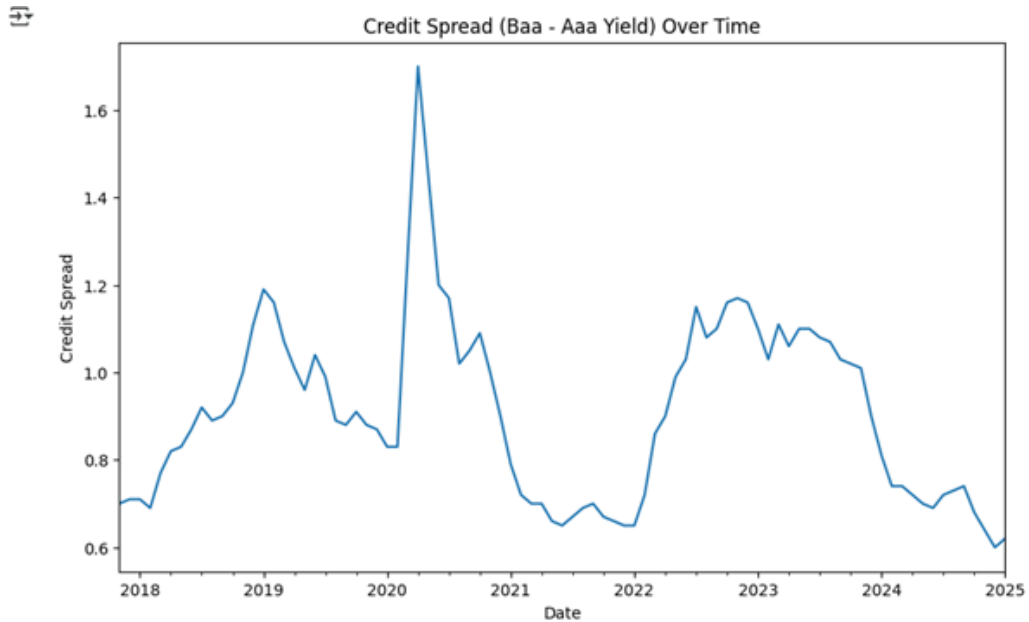


Fig. 14. Credit Spread

The line chart demonstrates the credit spread (Baa—Aaa yield) over time during financial turmoil. A steep rise was observed at the beginning of the COVID-19 epidemic, reflecting the heightened level of risk involved in lower-grade corporate bonds. After that, the spread saw a gradual decline, with variations that aligned with the changing market and the mood of the investors.

Distribution of Target Variable




Fig. 15. Target Variable Distribution

The bar chart describes the distribution of market direction, which is a target variable. A value of 1 indicates the market moved up, while 0 indicates downtime. The chart shows a slight imbalance; there were more instances of upward movement than downward days, which could affect the classification model's performance.

Confusion Matrix




Fig. 16. Confusion Matrix

The confusion matrix illustrated clearly portrays the neural network's executions in parallel with the predicted and actual labeling. The model rightly pointed out four downs (0) and five ups (1) of the market while wrongly classifying 1 of the upward movements as a downward and 2 of the downward ones as an upward movement. The fact that they have been relatively balanced in accuracy suggests that the model is reasonably operational, yet there remains the possibility of growth.

Journal Paper

The selection of the article (Wilamowski), "Neural Network Architectures and Learning Algorithms." The three key problems in the design of neural networks, namely, the selection of appropriate basic architecture, the selection of the right size for the network, and the choice of appropriate learning algorithms, are dealt with in his work. While the Multilayer Perceptron (MLP) is most frequently utilized, the author asserts it is outdated and often fails to deliver the same level of performance as sophisticated structures like Bridged MLP (BMLP). Another issue that Wilamowski points out is the error backpropagation (EBP) wherein he argues its inefficiency and unclear convergence are the reasons for its mass applicability. Conversely, it is demonstrated that training algorithms such as Levenberg-Marquardt (LM) and Neuron-by-Neuron (NBN) can drastically shorten the training time and harness performance ripples. These findings are particularly critical in financing, whereby the most efficient and quickest learning would be an advantage for modeling nonlinear and high-frequency time series.

Step 4: Technical Section

Hyperparameter Tuning for LDA, SVM, and Neural Network Models in Finance.

Linear Discriminant Analysis (LDA). Hyperparameter tuning is essential for the optimization of its performance in financial applications. While it is true that LDA is a more interpretable model, that is dependent on the accurate selection of its hyperparameters in order to strike a balance between bias and variance. Hyperparameters are set before training and influence the model's learning process.

General hyperparameter tuning methods:

- **Grid Search:** Systematically explores all possible combinations of hyperparameters from a set list to find the optimal configuration.
- **Cross-Validation:** Splits the data into several training and validation sets to test how well the model performs and adapts.
- **Regularization Techniques:** Addresses overfitting by adding penalty terms to the model's objective function.

Hyperparameter	Description	Example in Use
Solver	Algorithm used for optimization.	Used: 'eigen'. 'svd' was not used because it does not explicitly estimate the covariance matrix, which is required for shrinkage.
Shrinkage/Regularization Parameter	Helps when the within-class scatter matrix is singular or poorly conditioned. Applies regularization to the within-class scatter matrix to improve stability.	Used: 'auto' (Ledoit-Wolf shrinkage estimation).
Number of Components	Controls the number of discriminant axes used for dimensionality reduction.	Not explicitly set. For binary classification, LDA defaults to one component.
Priors	Allows setting class priors and influencing decision boundaries, especially with imbalanced classes.	Default/not set (estimated from training data distribution).

Tuning Techniques Used:

The LDA model's setup involved setting the solver and shrinkage parameters to 'eigen' and 'auto'. We chose these settings based on the dataset's features and our need to estimate shrinkage. We didn't use methods like GridSearchCV to fine-tune hyperparameters at this stage. However, picking 'eigen' and 'auto' was a conscious decision about hyperparameter values. To improve this further, we could look at more solver algorithms and shrinkage parameters. We might use techniques like GridSearch or Bayesian Optimization to do this.

SVMs. The SVM model is developed in such a manner that uses grid search (GridSearchCV) to determine the best estimator, regularization parameter C , kernel function (linear, poly, rbf etc.), gamma γ , epsilon ϵ hyperparameters.

Tuning Techniques Used: In the Jupyter notebook example:

- **For the classification task:** Upon grid searching, the following results obtained are as follows:
best kernel estimator: **.SVC**; C : **0.01**; kernel: **"poly"**,
- **For the regression task,** we selected the kernel: **"rbf"**, and proceeded to apply the grid search to determine the best parameters for the regularization parameter C , gamma γ , and epsilon ϵ . The results obtained are as follows:
 C : 100; γ : 0.01, ϵ : 0.0001
- Additionally, when performing grid search, parameters **probability** was set to **True** and **cv** (number of cross validation steps) was set as **5**

Neural Networks. Hyperparameter tuning in neural networks is an essential factor responsible for improving predictive performance, particularly in finance, where the models need to respond to complex nonlinear issues due to the involvement of volatile and interdependent variables. In contrast to parameters like weights and biases, hyperparameters are set before training starts and determine the network's learning behavior, architecture, and generalization ability.

Hyperparameter	Description	Example in Use
Learning Rate (η)	Step size that is used during weight update and impacts convergence speed and stability.	Used: 0.001 with adam optimizer for steady convergence.
Number of Epochs	This is the total number of passes through training data. If it is too few, underfitting occurs, and if it is too many, overfitting occurs as well.	Used: 300 for MLPClassifier.
Batch Size	This is the number of samples per training iteration, and it affects convergence and generalization	Default mini-batch used by MLPClassifier.
Hidden Layers / Neurons	This controls the model complexity. If we have more layers, they capture more patterns.	Used: 2 layers with (16, 8) neurons (MLP).
Activation Function	This introduces non-linearity into the network.	Used: ReLU for hidden layers, Sigmoid for output (binary class).
Dropout Rate	This randomly disables neurons during training to reduce overfitting.	Not applied in scikit-learn's MLP, but often 0.3–0.5 in deep learning.

Optimizer	This is an algorithm that is used for weight updates.	Used: Adam, a robust optimizer in financial models.
Weight Initialization	This is the initial setting of neuron weights. If we have poor values, they can delay convergence.	Automatically initialized by scikit-learn.
Regularization	This adds a penalty to loss to reduce overfitting.	L2 regularization applied internally in MLPClassifier.

Tuning Techniques Used:

- Manual tuning during experimentation is based on classification report metrics themselves.
- Cross-validation split was not used because of its time-series nature (shuffle=False).
- Furthermore, tuning can be conducted using GridSearchCV, Bayesian Optimization, or Keras Tuner in deep learning frameworks like TensorFlow/Keras.

Step 5: Marketing Alpha

Linear Discriminant Analysis.

Its adoption in finance has provided a structured and interpretable method for classification-based decision-making. It can maximize how well classes are separated while cutting down dimensions, making it a key tool to classify risks, spot market patterns, and score credit. By using statistical properties, LDA brings efficiency and clarity to predictive modeling, which is crucial for financial professionals seeking transparency and robustness in their models.

Key Strengths of LDA in Financial Applications:

- **Financial Market Classification:** LDA effectively differentiates between bullish and bearish market regimes using macroeconomic and market indicators.
- **Risk Categorization:** It enhances credit scoring and financial distress prediction. This happens by grouping borrowers or firms into set risk categories.
- **Computational Efficiency:** Unlike complex deep learning models, LDA is computationally lightweight, making it ideal for real-time decision-making.
- **Interpretability and Transparency:** Provides clear linear decision boundaries. These are key for following regulations and managing risk.
- **Dimensionality Reduction:** It projects high-dimensional financial data onto a lower-dimensional space. This keeps classification accuracy high and lowers the risk of overfitting.
- **Robustness with Small Datasets:** Performs reliably even when financial datasets are limited. This is different from machine learning models that need lots of data.

Real-World Use Cases:

Use Case	Description
Market Regime Classification	LDA identifies shifts in market conditions (e.g., bullish vs. bearish regimes) based on asset returns and macroeconomic indicators.
Credit Risk Assessment	Banks and lenders apply LDA to group borrowers by risk level when deciding on loans and their rates.
Portfolio Risk Analysis	It helps asset managers sort investments based on how returns are spread out and what's happening in the market.
Algorithmic Trading Signal Classification	Enhances trading models by classifying market states based on technical and fundamental indicators.

Strategic Edge: LDA's simplicity, efficiency, and transparency make it a valuable machine-learning model for financial applications. Unlike models that hide their workings (so-called "black-box" models), LDA shows how it makes decisions. This allows risk managers and traders to act with certainty. LDA can adapt to various areas of finance, from stocks to credit risk, which means it will stay useful in quantitative finance for a long time.

SVMs

Real-World Use Cases: SVM has a wide range of applications in finance.

1. **Financial Forecasting:** The ability of SVM models to work well with high dimensional data, minimise classification errors, and maximize margin classifier/regressor is a key strength required for detecting stock patterns and predicting future stock trends.
2. **Credit Risk Modelling:** SVM regressors are good predictors with minimal mean-squared errors, hence are better suited for prediction of likelihood of loan defaults of customer considering financial history: income, credit history, etc.
3. **Algorithmic Trading:** While SVMs are sensitive to noisy data points, but good for model predictors, they help to pick trading signals in the training set and identify trading opportunities in the test set. This could help algorithmic trading make profits from trades.

Strategic Edge:

1. SVMs perform well in high-dimensional spaces with many features, therefore suitable for macroeconomic and technical indicators in the stock market.
2. SVMs are robust to overfitting and hence help detect data outliers (i.e. anomalies) and are useful for fraud detection.
3. SVMs work well with small-scale data and can be useful for strategic financial investment decision-making in cases where there is quite a number of features to select from, however limited in size.

4. SVMs results are interpretable; as such, they are suited better for financial demand interpretation than neural networks.

Neural Networks.

Its introduction in the financial industry has provided a radical shift in the method of processing intricate, high-dimensional, and nonlinear data. Because of their flexibility and adaptability, they become strategic tools for producing alpha in portfolio management, detecting trading signals, and forecasting macroeconomic trends.

Key Strengths of Neural Networks in Financial Applications:

- Financial data is typically non-linear. Neural networks can effectively learn non-linear relationships between inputs and outputs, making them a perfect tool for modeling asset price behavior, volatility, and market regimes.
- Neural networks manage considerable feature sets, allowing the easy integration of technical indicators, sentiment ratings, macroeconomic elements, and diverse alternative data.
- Neural networks can stay up to date even when market conditions change because they can generalize from the data, and they can also retrain themselves.
- Markets are undeniably chaotic and unpredictable. Neural networks can tolerate incomplete, inaccurate, or partially damaged data because they are built on two specific principles: redundancy and regularization. As a result, they tend to work quite well under such circumstances.
- New sectors can be targeted using pre-trained financial models, which denote time efficiency and intrigue accuracy power. Models, from stocks to cryptocurrencies, are quite versatile and can be configured for different sectors.

Real-World Use Cases:

Use Case	Description
Stock Movement Prediction	We can forecast daily directional moves of the S&P 500 or BTC based solely on macro and market data.
Credit Risk Assessment	We can predict the likelihood of loan default by using borrower data and macro variables.
Portfolio Optimization	We can determine optimal asset weights by using pattern recognition in historical returns.
Sentiment-Driven Trading	We can use text-based financial news and social media processed through RNN/CNN models.

Strategic Edge: Neural networks in organizations can make more precise risk predictions, detect new trends as early as possible, and outperform other adaptive learning systems. These models' scalability, speed, and customizability make them robust machine-learning assets for dynamic markets.

Step 6: Learn More

- Akhter, Nahid. "Linear Discriminant Analysis in the Financial Market." *DataDrivenInvestor*, 20 May 2019, www.datadriveninvestor.com/2019/05/20/linear-discriminant-analysis-in-the-financial-market/.
 - **Description:** This article discusses the application of Linear Discriminant Analysis (LDA) to classify market movements and predict financial trends. Akhter explores how LDA can be used to analyze stock price behavior, helping investors make informed decisions by distinguishing between various market conditions. The article highlights LDA's potential in portfolio management and financial forecasting.
- Alrasheedi, Melfi. "Predicting Up/down Direction Using Linear Discriminant Analysis and Logit Model: The Case of SABIC Price Index." *Research Journal of Business Management*, vol. 6, no. 4, Apr. 2012, pp. 121–33, <https://doi.org/10.3923/rjbm.2012.121.133>.
 - **Description:** This paper uses Linear Discriminant Analysis (LDA) and the Logit model to predict the direction of price movements for the SABIC Price Index. Alrasheedi demonstrates how these statistical methods, combined with both technical and fundamental data, can enhance market trend forecasting, offering useful tools for investors seeking to classify upward and downward price movements.
- Altman, Edward I. "Financial Ratios, Discriminant Analysis and the Prediction of Corporate Bankruptcy." *The Journal of Finance*, vol. 23, no. 4, Sept. 1968, pp. 589–609, <https://doi.org/10.1111/j.1540-6261.1968.tb00843.x>.
 - **Description:** A foundational paper in financial risk modeling, this study applies discriminant analysis to predict corporate bankruptcy using financial ratios. Altman introduces the Z-score model, demonstrating how a weighted combination of key financial metrics can effectively classify firms into solvent and insolvent categories. This research has since served as the basis for numerous credit risk assessment models and remains widely used in financial distress prediction.
- I O Eweoya *et al* 2019 , "Fraud Detection using Support Vector Machines." *J. Phys.: Conf. Ser.* 1299 012039, DOI 10.1088/1742-6596/1299/1/012039
 - **Description:** This article clearly outlines the benefit of applying machine learning models like Support Vector Machine in detecting possible fraud in loan applications through hidden trends of data, which could result in a possible loan default. Results from the article show high accuracy for SVMs, as observed in the high number of true positives and low number of false positives. However, it proposes the integration of ensemble methods to improve model performance.
- Madhu, B. , Rahman, M. , Mukherjee, A. , Islam, M. , Roy, R. and Ali, L. (2021) "A Comparative Study of Support Vector Machine and Artificial Neural Network for Option Price Prediction". *Journal of Computer and Communications*, 9, 78-91. doi: 10.4236/jcc.2021.95006.

- **Description:** This paper compares the model performance of the ANN and SVM in option price prediction. Both models are tested with a publicly available dataset, namely SPY option price-2015, in both testing and training phases. The converted data through Principal Component Analysis (PCA) is used in both models to achieve better prediction accuracy. The RMSE value for both models was used to evaluate the performance of both models. Interestingly, the ANN was a better model predictor than SVM as it has a lower RMSE value.
- Mahmoud K. Okasha "Using Support Vector Machines in Financial Time Series Forecasting." Department of Applied Statistics, Al-Azhar University, Gaza Palestine, International Journal of Statistics and Applications, p-ISSN: 2168-5193 e-ISSN: 2168-5215, 2014; 4(1): 28-39, doi:10.5923/j.statistics.20140401.03
 - **Description:** This article describes the application of support vector machines in stock financial times series forecasting. The complexity of stock price indices and the noise of the time series further complicate forecasting. Hence, in this paper, algorithms ARIMA, ANN, as well as SVM, were fitted to Al-Quds Palestinian Stock Exchange ultimately to forecast two-month further data. The performance of all three models was compared using root-mean-squared error. It was concluded from the paper that SVM was a more accurate model for forecasting.
- Moghaddam, Amin Hedayati, et al. 'Stock Market Index Prediction Using Artificial Neural Network'. *Journal of Economics, Finance and Administrative Science*, vol. 21, no. 41, Dec. 2016, pp. 89–93. *ScienceDirect*, <https://doi.org/10.1016/j.jefas.2016.07.002>.
 - **Description:** This research assesses the proficiency of feedforward artificial neural networks in projecting the NASDAQ index utilizing short-term historical prices and day-of-week data. A multitude of neural network frameworks and training algorithms were experimented with. The best configuration was the 20-40-20 framework, which was, in fact, the configuration setting of the best model. The model was trained on the OSS algorithm and eventually achieved an accuracy of $R^2 \approx 0.94$. The results support that neural networks have the potential to be nonlinear financial time series forecasting tools for short-term market prediction.
- Patel, Jigar, et al. 'Predicting Stock and Stock Price Index Movement Using Trend Deterministic Data Preparation and Machine Learning Techniques'. *Expert Systems with Applications*, vol. 42, no. 1, Jan. 2015, pp. 259–68. *ScienceDirect*, <https://doi.org/10.1016/j.eswa.2014.07.040>.
 - **Description:** The paper above discusses the performance of four machine learning methods popularly referred to as ANN, SVM, Random Forest, and Naive Bayes. Historical data from the Indian market has been used to predict stock indices' movements. Besides, we have introduced a new method of input that determines the trend by recognizing continuous technical indicators and turning them into discrete trend signals. The results demonstrated that this approach improved the accuracy of predictions, and Random Forest and ANN performed the best.

- Shin, Kyung-Shik, and Yong-Joo Lee. "A Genetic Algorithm Application in Bankruptcy Prediction Modeling." *Expert Systems with Applications*, vol. 23, no. 3, Oct. 2002, pp. 321–28, [https://doi.org/10.1016/s0957-4174\(02\)00051-9](https://doi.org/10.1016/s0957-4174(02)00051-9).
 - **Description:** Shin and Lee apply genetic algorithms to improve bankruptcy prediction models, contrasting this approach with traditional statistical methods. Their work demonstrates the effectiveness of genetic algorithms in optimizing predictive models, resulting in improved accuracy for assessing the likelihood of corporate financial distress, with potential applications in financial decision-making.
- Wilamowski, Bogdan. 'Neural Network Architectures and Learning Algorithms'. *Industrial Electronics Magazine, IEEE*, vol. 3, Jan. 2010, pp. 56–63. *ResearchGate*, <https://doi.org/10.1109/MIE.2009.934790>.
 - **Description:** This paper evaluates issues related to the design of neural networks, among which are the selection of adequate architectures, the number of the network's units, and the algorithms used for their training. The proportion of the multilayer perceptron (MLP) and the learning method called error back-propagation (EBP) in many fields is the main subject of this study, and the author expresses the opinion that these methods are not always the best. The author proposes that more sophisticated structures, such as the bridged MLP (BMLP), and training methods, such as Levenberg-Marquardt (LM) and Neuron-by-Neuron (NBN) be used, especially for multi-variant complex and high-dimensional problems such as financial forecasting.

Step 7: Comparing Models

Feature	LDA	SVM	Neural Networks
Handling of Missing Data	LDA doesn't deal with missing data; you need to preprocess by removing or filling in gaps. This can result in lost information or skewed distributions if you're not careful.	SVM model training does not perform well in the presence of missing (NaN) values. Often times we handle missing by replacing NaN values or removing features with a large number of NaN values during data preprocessing.	Preprocessing must be done before the training begins by using Imputation or forward-fill technology.
Sensitivity to Outliers	Sensitive to outliers, as financial anomalies can distort classification performance.	Outliers significantly affect the decision boundary and support vectors of SVMs. However, this attribute could be useful for spotting data anomalies.	Without dropout or batch normalization, it can be robust with regularization and dramatic without regularization.

Computational Complexity	Computationally efficient, faster than complex non-linear classifiers like SVMs and NNs.	The computational efficiency of SVMs relies on the type of kernel function, number of samples, and features. For example. Linear kernel functions work well with large samples but fewer features. Reverse is the case for non-linear functions.	Varies widely based on network size and depth. Generally high, deep networks require substantially more computation and tuning.
Interpretability	Provides clear decision boundaries, which helps explain things for regulation and risk management.	SVMs are more interpretable than neural networks if probabilistic parameters are clearly defined.	Very low, especially for “black box” models, but can be enhanced with SHAP/LIME.
Performance with High - Dimensional Data	Works well with lots of data points while keeping important financial insights.	SVMs perform well in high-dimensional spaces with many features, therefore suitable macroeconomic and technical indicators in the stock market. Sometimes they perform better than neural networks like ANN.	Excellent, considering that it combines macro, market, and technical indicators. Requires careful architecture design.
Ability to Model Non-Linear Relationships	Assumes straight-line decision boundaries and has trouble with complex non-linear financial classification problems.	Perfect for non-linear relationships, as kernel transformations are often times required to obtain clear decision boundaries.	Exceptional as it can approximate a continuous function when it is given a sufficient depth.
Hyperparameter Tuning Complexity	Involves solver selection, shrinkage/regularization, number of components, and priors.	Hyperparameter tuning is simple but could be complicated when working with certain kernels.	Very high and complex, and involves many hyperparameters like layers, units, and learning rates that go together.

Required Data Preprocessing	Assumes normally distributed predictors, which is somewhat reasonable for financial metrics after transformation. Daily returns calculation is needed for stationarity.	SVM Classification and Regression tasks require data preprocessing primarily to remove outliers and irrelevant labels.	Here, input scaling (like standard scaler) is necessary for efficient training.
Overfitting Risk	Overfitting risk is relatively low compared to more complex models.	Overfitting risk is moderate as this can be controlled by the regulation parameter C.	Moderate to high because it needs dropout, early stopping, weight decay, or L2 regularization
Training Time	Training time is trains faster than SVM and NN.	Training time is pretty fast mainly for non-linear functins but different for linear kernel functions with many features.	Moderate to extremely high when working with larger networks or high-frequency data. Often requiring specialized hardware like GPUs.